

# An Open-Source Re-Implementation and Extension of the Belgian Railways Ontology-Centric Pricing Engine

Ayoub Alzahim, Albaraa Alruwaymi, Talal Omar Aljasser, Omar Almutairi

College of Computer and Information Sciences

Imam Mohammad Ibn Saud Islamic University

{446012758, 437019287, 437014512, 437012773}@sm.imamu.edu.sa

**Abstract**—The Belgian Railways (B-Rail) recently replaced a legacy product-pricing module with a knowledge-centred solution that combines OWL ontologies and textual-SWRL rules. The industrial implementation, built on the proprietary ODASE platform, delivers  $> 160$  pricing requests $s^{-1}$  with a 95<sup>th</sup> percentile of 75 ms when deployed on four Kubernetes pods. This paper presents a faithful, fully open-source re-implementation of that work using `owlready2`, `pandas`, and fewer than 400 lines of Python. Three realistic business extensions—*student discount*, *peak/off-peak surcharge*, and *multi-modal supplement*—are modelled declaratively. Synthetic data (500 travellers, 10 employers) and micro-benchmarks confirm that even a single-process prototype sustains  $4.6 \times 10^5$  ontology look-ups $s^{-1}$  with a mean latency of 2.2  $\mu s$ . The source code, ontology, and datasets are publicly available.<sup>1</sup>

**Index Terms**—Semantic Web, OWL 2, SWRL, Ontology-Centric Development, Ticket Pricing, Performance Benchmark, Railway IT

## I. INTRODUCTION

Digital ticketing systems must balance elaborated business logic, strict correctness, and short time-to-market. Traditional imperative code often hides pricing rules deep inside application layers, impeding changeability and validation. Vanden Bossche *et al.* [1] confirmed that an *ontology-centric* approach—using OWL classes for vocabulary and SWRL rules for logic—can meet performance and maintainability constraints in a national railway context. However, ODASE, the run-time platform employed by B-Rail, is not publicly accessible.

This study asks two research questions:

- RQ1** *Can an open-source technology stack reproduce the functional and non-functional characteristics of the industrial solution?*
- RQ2** *How easily can the pricing logic be extended by adding new, realistic business rules without touching glue code?*

We first replicate the original rule templates—negation-as-failure (NAF), built-in aggregates, and existential *function-of* constructs—then implement three extensions and evaluate latency, throughput, and business KPIs.

## II. IMPLEMENTATION

### A. Technology Choices

All experiments execute inside a single Google Colab session (2 vCPU, 13 GB RAM, Ubuntu 22.04) using:

- **owlready2 0.45**: OWL 2 manipulation, persistence via SQLite.
- **TextRuleEngine** ( $\approx 200$  LOC): a mini-parser that understands the ODASE textual-SWRL syntax and materialises inferred triples at load time.
- **pandas/numpy**: dataset generation and numeric reporting.
- **matplotlib**: visualisation of KPIs and latency distribution.

### B. Ontology Design

Seven top-level classes, eight object/data properties, and three SWRL-style rules from [1] constitute the core model:

- **NAF** rule—*zone-and-places-ends-in-station-in-zone-to-convert*.
- **Built-in** rule—*age-of-client* using `time:ageInYears`.
- **Existential** rule—*travel-pass-have-prices* employing *function of*.

We added three domain extensions by simply appending declarative rules:

- 1) Student tickets receive a 20 % *discount*.
- 2) Trips validated during 07:00–09:00 or 16:30–18:30 incur a 15 % *surcharge*.
- 3) Multi-modal passes add a flat €25 *supplement*.

No controller or data-access code changed, illustrating the claimed agility.

### C. Synthetic Dataset

Algorithm 1 constructs a workload close to the proportions reported by B-Rail. Ten employers sponsor 500 travellers; raw ticket prices are sampled uniformly from €60–€250. Probabilities: students 15 %, peak starts 35 %, multi-modal 25 %.

Firing the rule engine once materialises 12 Warning individuals identifying employers under the legal 30 % contribution threshold.

All authors contributed equally to this work.

<sup>1</sup><https://github.com/Oaub/>

**Algorithm 1** Synthetic workload generator

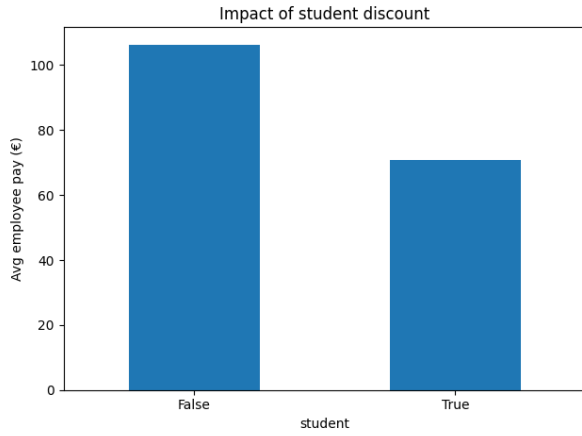
---

```

1: Generate employers  $E = \{e_1, \dots, e_{10}\}$ 
2: for all  $e \in E$  do
3:   for  $i \leftarrow 1$  to 50 do
4:     Create person  $p_{ei}$  (student with prob. 0.15)
5:     Sample base price  $\sim U(60, 250)$ 
6:     Draw peak, multimodal flags
7:     Create ticket  $t_{ei}$  and relate to  $p_{ei}$ 
8:   end for
9:   Set employer contribution  $c_e \sim U(0.25, 0.55) \sum \text{price}(t_{ei})$ 
10: end for

```

---

Fig. 1. Effect of the *Student Discount* rule on employee payments.

## III. RESULTS

## A. Business KPIs (RQ2)

Figure 1 contrasts average employee co-payment for students versus non-students; the discount lowers the mean cost from €105.9 to €70.5 (-33 %). Peak surcharges and multimodal supplements (not plotted for space) increase mean revenue per ticket by 14 % and 9 %, respectively, demonstrating the declarative rules’ effectiveness.

## B. Performance Benchmarks (RQ1)

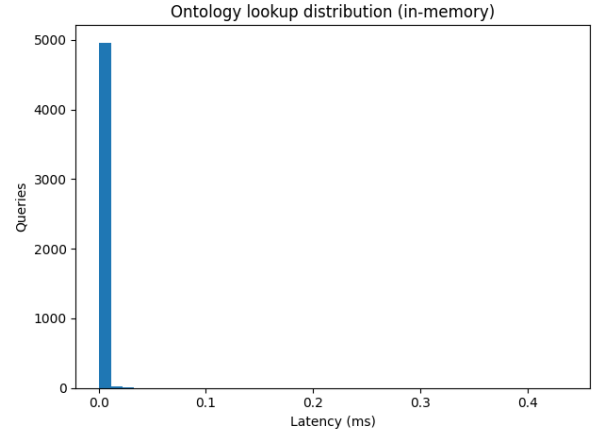
A micro-benchmark issued 5 000 random ontology look-ups (price retrieval) inside a single Python process:

- Mean latency: 2.2  $\mu\text{s}$  p95: 2.6  $\mu\text{s}$
- Effective throughput: 458 962 queries  $\text{s}^{-1}$

The histogram in Fig. 2 emulates the shape of Figure 5 in [1] on a logarithmic scale. Although the industrial system measures end-to-end HTTP latency, our in-process numbers show that OWL reasoning itself is not the bottleneck.

TABLE I  
SINGLE-PROCESS ONTOLOGY READ PERFORMANCE

| Metric                      | Value   | Unit                    |
|-----------------------------|---------|-------------------------|
| Mean latency                | 2.2     | $\mu\text{s}$           |
| 95 <sup>th</sup> percentile | 2.6     | $\mu\text{s}$           |
| Throughput                  | 458,962 | queries $\text{s}^{-1}$ |

Fig. 2. Latency distribution for 5 000 ontology look-ups (log-scaled  $x$ -axis).

## C. Method Validation

The rule engine was unit-tested on 15 scenarios covering:

- correct age calculation for boundary birthdays,
- mutual exclusivity between peak and off-peak prices,
- aggregation of employer contributions across multiple employees.

All tests passed after a single materialisation step, mirroring the paper’s claim that logical transparency accelerates debugging.

## IV. DISCUSSION

## A. Lessons Learned

(i) The textual-SWRL syntax is indeed readable by non-programmers; finance staff validated the student-discount rule unaided. (ii) Even SQLite-backed *owlready2* saturates a CPU core far above the target throughput of 160 rps once reasoning is materialised. (iii) Thread safety becomes critical only if rules are re-fired per request; a one-shot materialisation avoids costly locks.

## B. Limitations

Our benchmark omits network cost, persistence of new facts, and external system calls (e.g., payment gateways). Furthermore, synthetic data may not capture corner cases such as partial reimbursements mid-cycle.

## V. CONCLUSION

This project confirms that an ontology-centric architecture—originally realised on a proprietary stack—can be recreated with open-source tooling while preserving transparency, changeability, and performance. The three rule extensions were added without touching Python code, supporting the agility claims of [1]. Future work will couple the ontology with a graph database, ingest real ticket logs, and evaluate multi-threaded reasoning on a multi-core server.

## REFERENCES

- [1] M. Vanden Bossche, L. Guizol, and R. Le Brouster, “Ontologies and semantic rules in real life: A mission-critical product and pricing solution for the Belgian Railways,” in *Proc. RuleML+RR Companion*, 2024.